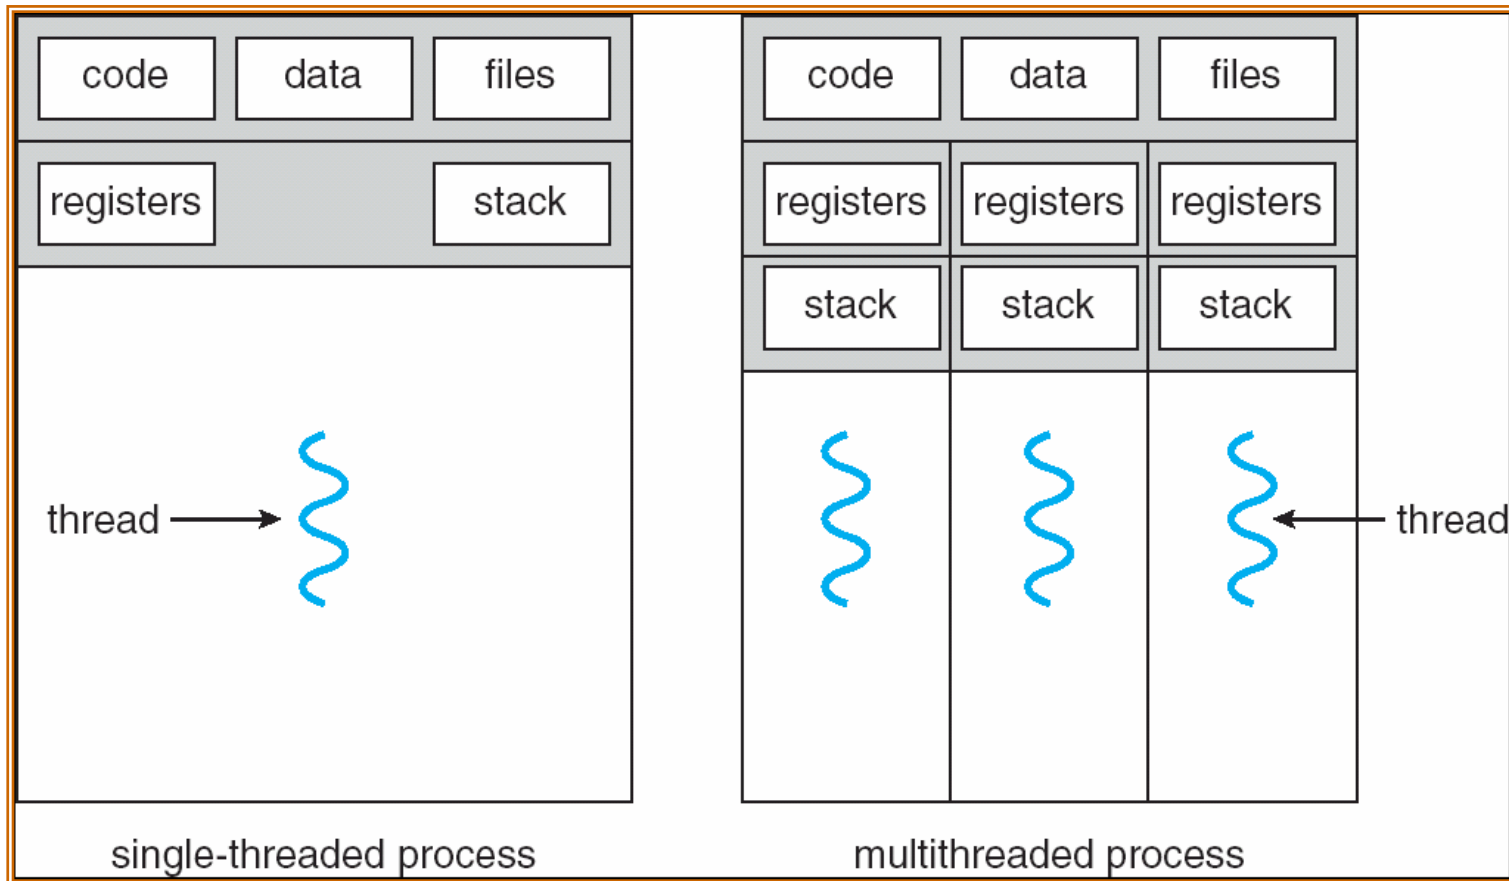


An Introduction to Programming with Threads

- Read the Birrell paper
 - excellent introductory paper
 - promotes understanding the material
 - abstract content with direct application
 - limited and rather outdated concrete technical content

Threads

- a thread is a single sequential flow of control
 - a process can have many threads and a single address space
 - threads share memory and, hence, need to cooperate to produce correct results
 - thread has thread specific data (registers, stack pointer, program counter...)



Why use threads

- Threads are useful because of real-world parallelism:
 - input/output devices (flesh or silicon) may be slow but are independent
 - distributed systems have many computing entities
 - multi-processors are becoming more common, and future platforms will all be multi-core
 - better resource sharing & utilization than processes

Thread Mechanisms

- Birrell identifies four mechanisms used in threading systems:
 - thread creation
 - mutual exclusion
 - waiting for events
 - interrupting a thread's wait
- In most mechanisms in current use, only the first three are covered
- primitives used abstract, not derived from actual threading system or programming language!

Example Thread Primitives

- Thread creation
 - Thread type
 - `Fork(proc, args)` returns thread
 - `Join(thread)` returns value
- Mutual Exclusion
 - Mutex type
 - `Lock(mutex)`, a block-structured language construct in this lecture

Example Thread Primitives

- Condition Variables
 - `Condition` type
 - `Wait(mutex, condition)`
 - `Signal(condition)`
 - `Broadcast(condition)`
- Fork, Wait, Signal, etc. are not to be confused with the UNIX “fork”, “wait”, “signal”, etc. calls

Creation Example

```
Thread thread1;  
thread1 = Fork(safe_insert, 4);  
safe_insert(6);  
Join(thread1); // Optional
```


Mutex Example

```
list<int> my_list;
```

```
Mutex m;
```

```
void safe_insert(int i) {  
    Lock(m) {  
        my_list.insert(i);  
    }  
}
```

Condition Variables

- Mutexes are used to control access to shared data
 - only one thread can execute inside a Lock clause
 - other threads who try to Lock, are blocked until the mutex is unlocked
- Condition variables are used to wait for specific events
 - free memory is getting low, wake up the garbage collector thread
 - 10,000 clock ticks have elapsed, update that window
 - new data arrived in the I/O port, process it
- Could we do the same with mutexes?
 - (think about it and we'll get back to it)

Condition Variable Example

```
Mutex io_mutex;
```

```
Condition non_empty;
```

```
...
```

```
Consumer:
```

```
Lock (io_mutex) {  
    while (port.empty())  
        Wait(io_mutex, non_empty);  
    process_data(port.first_in());  
}
```

```
Producer:
```

```
Lock (io_mutex) {  
    port.add_data();  
    Signal(non_empty);  
}
```

Condition Variables Semantics

- Each condition variable is associated with a single mutex
- *Wait atomically* unlocks the mutex and blocks the thread
- **Signal** awakes a blocked thread
 - the thread is awoken inside *Wait*
 - tries to lock the mutex
 - when it (finally) succeeds, it returns from the *Wait*
- Doesn't this sound complex? Why do we do it?
 - the idea is that the “condition” of the condition variable depends on data protected by the mutex

Condition Variable Example

```
Mutex io_mutex;
```

```
Condition non_empty;
```

```
...
```

```
Consumer:
```

```
Lock (io_mutex) {  
    while (port.empty())  
        Wait(io_mutex, non_empty);  
    process_data(port.first_in());  
}
```

```
Producer:
```

```
Lock (io_mutex) {  
    port.add_data();  
    Signal(non_empty);  
}
```

Couldn't We Do the Same with Plain Communication?

```
Mutex io_mutex;  
...  
Consumer:  
Lock (io_mutex) {  
    while (port.empty())  
        go_to_sleep(non_empty);  
    process_data(port.first_in());  
}  
Producer:  
Lock (io_mutex) {  
    port.add_data();  
    wake_up(non_empty);  
}
```

- What's wrong with this? What if we don't lock the mutex (or unlock it before going to sleep)?

Mutexes and Condition Variables

- Mutexes and condition variables serve different purposes
 - Mutex: exclusive access
 - Condition variable: long waits
- *Question:* Isn't it weird to have both mutexes and condition variables? Couldn't a single mechanism suffice?
- *Answer:*

Use of Mutexes and Condition Variables

- Protect shared mutable data:

```
void insert(int i) {  
    Element *e = new Element(i);  
    e->next = head;  
    head = e;  
}
```

- What happens if this code is run in two different threads with no mutual exclusion?

Using Condition Variables

```
Mutex io_mutex;
Condition non_empty;
...
Consumer:
Lock (io_mutex) {
    while (port.empty())
        Wait(io_mutex, non_empty);
    process_data(port.first_in());
}
Producer:
Lock (io_mutex) {
    port.add_data();
    Signal(non_empty);
}
```

Why use **while** instead of **if**? (think of many consumers, simplicity of coding producer)

Readers/Writers Locking

```
Mutex counter_mutex;
Condition read_phase,
           write_phase;
int readers = 0;
Reader:
Lock(counter_mutex) {
    while (readers == -1)
        Wait(counter_mutex,
             read_phase);
    readers++;
}
... //read data
Lock(counter_mutex) {
    readers--;
    if (readers == 0)
        Signal(write_phase);
}
```

```
Writer:
Lock(counter_mutex) {
    while (readers != 0)
        Wait(counter_mutex,
             write_phase);
    readers = -1;
}
... //write data
Lock(counter_mutex) {
    readers = 0;
    Broadcast(read_phase);
    Signal(write_phase);
}
```

Comments on Readers/Writers Example

- Invariant: $\text{readers} \geq -1$
- Note the use of Broadcast
- The example could be simplified by using a single condition variable for phase changes
 - less efficient, easier to get wrong
- Note that a writer signals all potential readers and one potential writer. Not all can proceed, however
 - (*spurious wake-ups*)
- Unnecessary lock conflicts may arise (especially for multiprocessors):
 - both readers and writers signal condition variables while still holding the corresponding mutexes
 - Broadcast wakes up many readers that will contend for a mutex

Readers/Writers Example

Reader:

```
Lock(mutex) {  
    while (writer)  
        Wait(mutex, read_phase)  
    readers++;  
}
```

... // read data

```
Lock(mutex) {  
    readers--;  
    if (readers == 0)  
        Signal(write_phase);  
}
```

Writer:

```
Lock(mutex) {  
    while (readers != 0 || writer)  
        Wait(mutex, write_phase)  
    writer = true;  
}
```

... // write data

```
Lock(mutex) {  
    writer = false;  
    Broadcast(read_phase);  
    Signal(write_phase);  
}
```

Avoiding Unnecessary Wake-ups

```
Mutex counter_mutex;  
Condition read_phase, write_phase;  
int readers = 0, waiting_readers = 0;
```

Reader:

```
Lock(counter_mutex) {  
    waiting_readers++;  
    while (readers == -1)  
        Wait(counter_mutex,  
              read_phase);  
    waiting_readers--;  
    readers++;  
}  
... //read data  
Lock(counter_mutex) {  
    readers--;  
    if (readers == 0)  
        Signal(write_phase);  
}
```

Writer:

```
Lock(counter_mutex) {  
    while (readers != 0)  
        Wait(counter_mutex,  
              write_phase);  
    readers = -1;  
}  
... //write data  
Lock(counter_mutex) {  
    readers = 0;  
    if (waiting_readers > 0)  
        Broadcast(read_phase);  
    else  
        Signal(write_phase);  
}
```

Problems With This Solution

- Explicit scheduling: readers always have priority
 - may lead to starvation (if there are always readers)
 - fix: make the scheduling protocol more complicated than it is now

To Do:

- Think about avoiding the problem of waking up readers that will contend for a single mutex if executed on multiple processors

Deadlocks (brief)

- We'll talk more later... for now beware of deadlocks
- Examples:
 - A locks M1, B locks M2, A blocks on M2, B blocks on M1
 - Similar examples with condition variables and mutexes
- Techniques for avoiding deadlocks:
 - Fine grained locking
 - Two-phase locking: acquire all the locks you'll ever need up front, release all locks if you fail to acquire any one
 - very good technique for some applications, but generally too restrictive
 - Order locks and acquire them in order (e.g., all threads first acquire M1, then M2)

- The scope of multithreading

LWPs, kernel and user threads

- kernel-level threads – supported by the kernel
 - Solaris, Linux, Windows XP/2000
 - all scheduling, synchronization, thread structures maintained in kernel
 - could write apps using kernel threads, but would have to go to kernel for everything
- user-level threads – supported by a user-level library
 - Pthreads, Java threads, Win32...
 - sched. & synch can often be done fully in user space; kernel doesn't need to know there are many user threads
 - problem with blocking on a system call


- LightWeight Processes - LWP
 - these are “virtual CPUs”, can be multiple per process
 - the scheduler of a threads library schedules user-level threads to these virtual CPUs
 - kernel threads implement LWPs => visible to the kernel, and can be scheduled
 - sometimes LWP & kernel threads used interchangeably, but there can be kernel threads without LWPs

Multithreading models

- There are three dominant models for thread libraries, each with its own trade-offs
 - many threads on one LWP (many-to-one)
 - one thread per LWP (one-to-one)
 - many threads on many LWPs (many-to-many)
- similar models can apply on scheduling kernel threads to real CPUs


Many-to-one

- In this model, the library maps all threads to a single lightweight process
- Advantages:
 - totally portable
 - easy to do with few systems dependencies
- Disadvantages:
 - cannot take advantage of parallelism
 - may have to block for synchronous I/O
 - there is a clever technique for avoiding it
- Mainly used in language systems, portable libraries





One-to-one

- In this model, the library maps each thread to a different lightweight process
- Advantages:
 - can exploit parallelism, blocking system calls
- Disadvantages:
 - thread creation involves LWP creation
 - each thread takes up kernel resources
 - limiting the number of total threads
- Used in LinuxThreads and other systems where LWP creation is not too expensive



Many-to-many

- In this model, the library has two kinds of threads: *bound* and *unbound*
 - bound threads are mapped each to a single lightweight process
 - unbound threads *may* be mapped to the same LWP
- Probably the best of both worlds
- Used in the Solaris implementation of Pthreads (and several other Unix implementations)



High-Level Program Structure Ideas

- Boss/workers model
- Pipeline model
- Up-calls
- Keeping shared information consistent using version stamps

Thread Design Patterns

Common ways of structuring programs using threads

- Boss/workers model
 - boss gets assignments, dispatches tasks to workers
 - variants (thread pool, single thread per connection...)
- Pipeline model
 - do some work, pass partial result to next thread
- Up-calls
 - fast control flow transfer for layered systems
- Version stamps
 - technique for keeping information consistent

Boss/Workers

Boss:

```
forever {  
    get a request  
    switch(request)  
        case X: Fork (taskX)  
        case Y: Fork (taskY)  
        ...  
}
```

Worker:

```
taskX();
```

- Advantage: simplicity
- Disadvantage: bound on number of workers, overhead of threads creation, contention if requests have interdependencies
- Variants: fixed thread pool (aka *workpile*, *workqueue*), producer/consumer relationship, workers determine what needs to be performed...

Pipeline

- Each thread completes portion of a task, and passes results
- like an assembly line or a processor pipeline
- Advantages: trivial synchronization, simplicity
- Disadvantages: limits degree of parallelism, throughput driven by slowest stage, handtuning needed

Up-calls

- Layered applications, e.g. network protocol stacks have top-down and bottom-up flows
- Up-calls is a technique in which you structure layers so that they can expect calls from below
- Thread pool of specialized threads in each layer
 - essentially an up-call pipeline per connection
- Advantages: best when used with fast, synchronous control flow transfer mechanisms or program structuring tool
- Disadvantages: programming becomes more complicated, synchronization required for top-down

Version Stamps

- (not a programming structure idea but useful technique for any kind of distributed environment)
- maintain “version number” for shared data
 - keep local cached copy of data
 - check versions to determine if changed